

John Munro

Last update on June 23, 2017

jbmunro4@gmail.com • (813) 546-0723 • www.linkedin.com/in/jbmunro4
100 Van Ness Ave #2602 • San Francisco, CA • 94102

Summary

Data nerd with strong development skills, a proven history of designing and building creative data products, an intuition for data architecture, and strong leadership capabilities.

I enjoy an intensely challenging and creative career in various forms of Information Science that deepen my existing knowledge of Machine Learning, Data Analysis, Artificial Intelligence, Systems Design, and

Cognitive Science while broadening my worldview. I want my views to be challenged on a regular basis and to always be learning new things.

I have a passion for understanding the nature of Intelligence and strive to formalize or leverage aspects of Human Intelligence in everything that I build, whether from an AI and Automation perspective or from the perspective of Product Design and Usability.

Experience

[eatsa](#)

SAN FRANCISCO, CA

Software Engineer #7 (Data Scientist & Data Engineer #1)

Mar '16 – May '17

Collaborated tightly with the analytics team to rapidly stand up the initial analytics infrastructure. Provided a centralized data warehouse to correlate and analyze information across a mix of proprietary and third party data sources across the engineering, marketing, product, and science teams. Worked closely with executive leadership and analytics team to provide frequent board updates and ad-hoc analysis.

Significantly improved accuracy and removed volatility in the customer-facing Order Delivery ETA provided to customers once they've purchased food through the mobile apps.

Technical Lead on project Blackbird, a very high priority initiative led by the CEO, to measure and improve operational efficiency. Rapidly designed and implemented a custom inventory management solution using google sheets after arduously working through complex operational requirements, accounting requirements, reporting requirements, and intensely dirty data. Generated product requirements and evaluated inventory management solutions, workforce management solutions, and Enterprise Resource Planning solutions (ERP's). Solely responsible for the vast majority of technical implementation for a highly customized ERP integration. This included eatsa's first production micro-service which was built in Python on top of the [serverless framework](#).

Data Science Lead on Menu Personalization and Dynamic Pricing. Collaborated closely with Product, Design, Engineering, the Food and Operations teams, the molecular Science team, Analytics, and executive leadership to design and implement the initial technical solutions for Menu Recommendations and Frequency-Based Discounts. Designed various pricing models in collaboration with Finance. Designed various menu recommendation algorithms in collaboration with the Science Team, Food Team, and Engineering. Designed and implemented Cibo, a production micro-service built on top of Django to serve personalized menus and prices.

[Anomali](#)

REDWOOD CITY, CA

Lead Data Scientist (Data Scientist #1)

May '15 – Mar '16

Individual contributor with a wide range of responsibilities across product development, design, and research. Built the next iteration of Anomali's machine learning-based scoring engine, worked on targeted research projects with security analysts, vetted new data sources, and contributed to the sales team by integrating Anomali's many data sources to give a 360° view of customers as they move through the sales process and use Anomali's SaaS security platform.

[Endgame](#)

SAN FRANCISCO, CA

Lead Data Scientist, Hiring Manager

Jun '14 – May '15

Built a diverse team of product-focused Data Scientists in an effort to create various analytical models for a variety of Cyber Security applications. Implemented cross-functional collaboration to evangelize our capabilities and learn more about potential use cases across the company. Initiated a culture of healthy peer review along with multiple successful initiatives to improve the team's security domain knowledge.

Designed and contributed to a data science platform that supports analytic R&D in both batch and streaming environments at scale. This platform services a new product aimed at giving System and

Security Administrators visibility into cloud based production environments to identify anomalous behavior indicative of security threats, misconfiguration, or inappropriate resource allocation.

WASHINGTON, DC

Senior Backend Developer, Tech Lead

Oct '12 – Jun '14

Researched, designed, and implemented a full stack product to identify trends in large, high dimensional data sets using unsupervised clustering methods to aid in data exploration. Iterated closely with the customer to define their needs and determine appropriate solutions which provide a significant advantage over their old methods and redefine how they think about their specific problems. Hired, onboarded, and led a small team of developers to support the product going forward.

ATLANTA, GA

Malware Analysis Engineer (Data Scientist #1)

Oct '11 – Oct '12

Researched and implemented a wide variety of statistical models to classify and categorize malicious characteristics of malware binaries, TCP packet headers, malware domain names, and Botnet Command and Control communication patterns. Presented [Clairvoyant Squirrel](#), a novel domain name classifier, at FloCon 2013 with [Jason Trost](#).

Autonomy (since purchased by HP)

ATLANTA, GA

Technology Specialist

Jan '11 – Oct '11

Worked as a presales engineer who frequently interacted with customers to understand their needs and communicate how they would benefit from Autonomy's solutions. Rapidly implemented full stack Proof of Concepts to demo Autonomy's capabilities with respect to customer needs. Frequently interacted with engineers to communicate customer requirements and debug solutions. Worked for three months on a customer site in the UK to help configure, deploy, and test new Autonomy products in a large production environment.

Please refer to my [Linked-in profile](#) for a more complete list of work experiences along with recommendations.

Education

Georgia Institute of Technology

ATLANTA, GA

Bachelor degree in Computer Science & Minor in Psychology

2006 – 2010

Specialized heavily in Artificial Intelligence, Cognitive Neuroscience and Computer Networking. Completed 3 years of Undergraduate Research focused on AI, Automated Story Telling, and Game Design with a thesis on computational models for automating machinima generation. Took courses in Machine Learning, Statistics, Experimental Design, Robotics, Computer Networking, Cognitive Psychology, and Neuropsychology.

Skills Overview

Applied and theoretical algorithmic expertise: Random Forests, KNN, SVM, Decision Trees, K-Means, K-Medoids, DBScan, Self Organizing Maps, Hierarchical Clustering, PCA, Information Gain Analysis, Linear Regression, and Holtz-Winters Forecasting.

Languages: Python, Bash, R, Javascript, HTML, CSS, Java (and some Scala recently for Spark)

Frameworks, tools, and libraries: UNIX, Django, numpy, scipy, scikit-learn, matplotlib, d3.js, jquery, PostgreSQL, Elasticsearch, s3, ec2, Redshift, HDFS, Map Reduce, Spark/PySpark, Storm, Pig, Looker, Tableau

Interests

Non-exhaustive and in alphabetical order: cooking, data visualization, eating, education, game design and development, robotics, snowboarding, UX and product design